American International University Bangladesh

Introduction to Data Science

Section: J

Group: J

Supervised by

KAMRUN NAHER KOLI

Name	ID	Contribution
SM FAHIM SHAHRI	23-50114-1	
MD RAFIT	23-50418-1	
MD.SIAMUZZAMAN SIAM	22-49113-3	

Table of Content

- **Introduction**
- **Dataset Description**
- **Implementing R program**
- **Introducing Data Issues**
- **Handling Missing Data**
- **Fixing Invalid Data Entries**
- **Duplicate Detection and Removal**
- **Fixing Date Format Issues**
- **Data Diagnostics**
- **Missingness Visualization**
- **Data Type Cleaning**
- **Handling Outliers**
- **Feature Engineering**
- **Normalization**
- **Handling Class Imbalance**
- **Exploratory Data Analysis (EDA)**
- **Compact Cleaning Script Overview**
- **Train–Test Split**
- **Creating Target Variable**
- **Conclusion**

Earthquake 1995-2023

Dataset Description:

This dataset contains 1,000 earthquake records with 19 attributes describing each event. It includes essential seismic information such as the earthquake's magnitude, depth, intensity scores (CDI and MMI), and tsunami indicators, along with geographic details like latitude, longitude, nearby location, country, and continent. Temporal data is captured through the date and time of each event, while additional technical fields like the number of reporting stations (nst), station distance (dmin), azimuthal gap, significance score, and magnitude type provide deeper analytical value. Some fields, such as alert level and continent, contain missing values. Overall, the dataset combines seismic activity metrics with geospatial and descriptive information suitable for analysis, prediction, and visualization of earthquake patterns.

Implementing R program and its output:

Code:

1. **Reading the csv file**

```
file_path <- "C:/Users/ASUS/Downloads/earthquack_project.csv"
df <- read.csv(file_path, stringsAsFactors = FALSE)


df <- readr::read_csv(file_path)
head(df, 15)
```

Output:

```
Rows: 1000 Columns: 19
── Column specification ────────────────────────────────────────────────────────
Delimiter: ","
chr  (8): title, date_time, alert, net, magType, location, continent, country
dbl (11): magnitude, cdi, mmi, tsunami, sig, nst, dmin, gap, depth, latitude, longitude

ℹ Use `spec()` to retrieve the full column specification for this data.
ℹ Specify the column types or set `show_col_types = FALSE` to quiet this message.
> head(df, 15)
# A tibble: 15 × 19
   title     magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType
   <chr>         <dbl> <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>
 1 M 6.5 -…        6.5 16-08-20…     7     4 green       0   657 us      114 7.18    25  mww
 2 M 6.5 -…        6.5 19-07-20…     8     6 yell…       0   775 us       92 0.679   40  mww
 3 M 6.6 -…        6.6 17-07-20…     7     5 green       0   899 us       70 1.63    28  mww
 4 M 7.2 -…        7.2 16-07-20…     6     6 green       1   860 us      173 0.907   36  mww
 5 M 7.3 -…        7.3 16-07-20…     0     5 NA          1   820 at       79 0.879  173. Mi
 6 M 6.6 -…        6.6 10-07-20…     5     4 green       1   802 us       95 2.45    37  mww
 7 M 6.9 -…        6.9 02-07-20…     4     4 green       1   741 us      136 1.18    23  mww
 8 M 7.2 -…        7.2 15-06-20…     8     6 green       1   804 us       85 2.59    24  mww
 9 M 6.6 -…        6.6 25-05-20…     6     6 green       1   733 us       50 2.16   129  mww
10 M 7.1 -…        7.1 20-05-20…     3     4 green       1   777 us       98 2.81    56  mww
11 M 7.7 -…        7.7 19-05-20…     5     4 green       1   927 us      277 3.11    15  mww
12 M 7.6 -…        7.6 10-05-20…     8     6 green       1   907 us      111 3.63    17  mww
13 M 6.6 -…        6.6 28-04-20…     0     2 green       0   670 us       52 5.12    50  mww
14 M 7.1 -…        7.1 24-04-20…     9     7 green       0   838 us       70 2.28    58  mww
15 M 7.1 -…        7.1 24-04-20…     8     7 green       1   777 us      126 0.708   23  mww
# ℹ 6 more variables: depth <dbl>, latitude <dbl>, longitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
```

Description:
Here the dataset was loaded using R in RStudio. The head(df, 15) output displays the first 15 instances of the earthquake dataset, showing all attributes such as magnitude, date, intensity levels, geographic location, and alert information. From this preview, it is clear that the dataset contains valid records along with some missing values, incomplete fields, and possible outliers in certain attributes. This initial view helps understand the structure and quality of the data before performing further analysis.

2. Introducing data issues

```
df_dirty <- df
```

1) Introduce missing values at random in magnitude and depth

```
n <- nrow(df_dirty)
idx_mag_na <- sample(1:n, size = round(0.03*n))   # 3% missing magnitudes
idx_depth_na <- sample(1:n, size = round(0.03*n)) # 3% missing depth
df_dirty$magnitude[idx_mag_na] <- NA
df_dirty$depth_km[idx_depth_na] <- NA
```

Output:

```
Error in `$<-`:
! Assigned data `<lgl>` must be compatible with existing data.
✕ Existing data has 1000 rows.
✕ Assigned data has 981 rows.
i Only vectors of size 1 are recycled.
Caused by error in `vectbl_recycle_rhs_rows()`:
! Can't recycle input of size 981 to size 1000.
Run `rlang::last_trace()` to see where the error occurred.
Warning message:
Unknown or uninitialised column: `depth_km`.
```

Solution of Introduce missing values at random in magnitude and depth

```
sum(is.na(df_clean$magnitude))
sum(is.na(df_clean$depth))

table(is.na(df_dirty$magnitude), is.na(df_clean$magnitude))
table(is.na(df_dirty$depth), is.na(df_clean$depth))
```

Output:

```
> sum(is.na(df_clean$magnitude))
[1] 0
> sum(is.na(df_clean$depth))
[1] 0
>
> table(is.na(df_dirty$magnitude), is.na(df_clean$magnitude))

        FALSE
  FALSE   970
  TRUE     30
> table(is.na(df_dirty$depth), is.na(df_clean$depth))

        FALSE
  FALSE  1000
> source("C:/Users/Rafit Tishan Neo/Downloads/Data Science Project/project.R")
>
```

Description:

The missing values introduced in the *magnitude* and *depth* columns were handled by applying median imputation. This method replaces each missing entry with the median of the corresponding attribute, ensuring that the overall distribution of the data remains stable and unaffected by outliers. This approach preserves dataset size, maintains statistical consistency, and prepares the data for further analysis without introducing bias.

Introduce some invalid magnitude values (strings, negatives)

2) Introduce some invalid magnitude values (strings, negatives)
bad_idx <- sample(1:n, size = round(0.01*n))
df_dirty$magnitude[bad_idx[1:floor(length(bad_idx)/2)]] <- "unknown"
df_dirty$magnitude[bad_idx[(floor(length(bad_idx)/2)+1):length(bad_idx)]] <- -9

Find string values ("unknown")

df_dirty[df_dirty$magnitude == "unknown",]

Output:

```
> df_dirty[df_dirty$magnitude == "unknown", ]
# A tibble: 35 × 19
   title         magnitude date_time   cdi   mmi alert tsunami   sig net     nst   dmin   gap magType depth latitude
   <chr>         <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl>  <dbl> <dbl> <chr>   <dbl>    <dbl>
 1 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
 2 M 7.8 - Cent… unknown   6/2/2023…     9     9 red         0  2910 us      118  1.92     32 mww      17.9     37.2
 3 M 6.6 - 55 k… unknown   21-04-20…     6     5 green       1   762 us        0  0.914    94 mww      27       11.6
 4 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
 5 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
 6 M 7.0 - 114 … unknown   17-07-20…     6     7 green       1   820 us        0  1.67     18 mww      73       -7.85
 7 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
 8 M 6.5 - 108 … unknown   15-08-20…     6     6 green       1   655 us        0  0.576    22 mww      33.9     51.4
 9 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
10 NA            NA        NA           NA    NA NA         NA    NA NA       NA NA        NA NA       NA       NA
# i 25 more rows
# i 4 more variables: longitude <dbl>, location <chr>, continent <chr>, country <chr>
# i Use `print(n = ...)` to see more rows
```

Find negative magnitudes

df_dirty[df_dirty$magnitude == -9,]

Output:

```
> df_dirty[df_dirty$magnitude == -9, ]
# A tibble: 35 × 19
   title magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType depth latitude longitude location continent
   <chr> <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>   <dbl>    <dbl>     <dbl> <chr>    <chr>
 1 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 2 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 3 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 4 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 5 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 6 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 7 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
 8 M 6.… -9        7/4/2016…     7     5 green       1   692 us        0  6.79    25 mww      27.6    -14.0      167. Sola, V… NA
 9 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
10 NA    NA        NA           NA    NA NA         NA    NA NA       NA NA       NA NA       NA       NA        NA  NA       NA
# i 25 more rows
# i 1 more variable: country <chr>
# i Use `print(n = ...)` to see more rows
>
```

Count how many invalid values exist

sum(df_dirty$magnitude == "unknown", na.rm = TRUE)

sum(df_dirty$magnitude == -9, na.rm = TRUE)

Output:

```
> sum(df_dirty$magnitude == "unknown", na.rm = TRUE)
[1] 5
> sum(df_dirty$magnitude == -9, na.rm = TRUE)
[1] 5
>
```

Description:

The magnitude field contained invalid values because approximately 1% of the dataset was intentionally corrupted by replacing half of the selected values with the string "unknown" and the other half with the negative number -9. Since earthquake magnitudes must be positive numerical values, these invalid entries disrupt statistical analysis and modeling.

Solutions

Convert magnitude to numeric

df_clean <- df_dirty

df_clean$magnitude <- suppressWarnings(as.numeric(df_clean$magnitude))

Replace negative values with NA

df_clean$magnitude[df_clean$magnitude < 0] <- NA

Impute missing/invalid values with the median
df_clean$magnitude[is.na(df_clean$magnitude)] <-

median(df_clean$magnitude, na.rm = TRUE)

Check if any invalid values remain

summary(df_clean$magnitude)

sum(df_clean$magnitude < 0, na.rm = TRUE)

sum(is.na(df_clean$magnitude))

Output:

```
> summary(df_clean$magnitude)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max. 
   6.50    6.60    6.80    6.93    7.10    9.10 
> sum(df_clean$magnitude < 0, na.rm = TRUE)
[1] 0
> sum(is.na(df_clean$magnitude))
[1] 0
>
```

Description:

To correct this issue, the magnitude column was first converted to numeric, which automatically converted all non-numeric entries to missing values. Negative magnitude values were also replaced with NA, as they do not represent real-world measurements. Finally, all missing or invalid magnitude values were imputed using the median of the valid magnitudes, ensuring that the dataset remains consistent and statistically reliable.

1) How many rows and duplicates?

nrow(df_dirty) # total rows

sum(duplicated(df_dirty)) # number of duplicate rows

head(df_dirty[duplicated(df_dirty),])

df_clean <- df_dirty[!duplicated(df_dirty),]

nrow(df_clean)

sum(duplicated(df_clean)) # should be 0

Output:

```
> # 1) How many rows and duplicates?
> nrow(df_dirty)                    # total rows
[1] 1000
> sum(duplicated(df_dirty))         # number of duplicate rows
[1] 0
>
> # 2) See some duplicate examples (complete-row duplicates)
> head(df_dirty[duplicated(df_dirty), ])
# A tibble: 0 × 19
# ℹ 19 variables: title <chr>, magnitude <chr>, date_time <chr>, cdi <dbl>, mmi <dbl>, alert <chr>, tsunami <dbl>, sig <dbl>,
#   net <chr>, nst <dbl>, dmin <dbl>, gap <dbl>, magType <chr>, depth <dbl>, latitude <dbl>, longitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
>
> # 3) Remove duplicates, keeping the first occurrence
> df_clean <- df_dirty[!duplicated(df_dirty), ]
>
> # 4) Check that duplicates are gone
> nrow(df_clean)                    # rows after cleaning
[1] 1000
> sum(duplicated(df_clean))         # should be 0
[1] 0
>
```

problem 4) Introduce inconsistent date formats in 'date_time' column

if("date_time" %in% names(df_dirty)){

samp <- sample(1:nrow(df_dirty), size = round(0.02 * nrow(df_dirty)))

keep only the date part from original, but in inconsistent format

random_dates <- Sys.Date() - sample(1:1000, length(samp), replace = TRUE)

df_dirty$date_time[samp] <- format(random_dates, "%d-%m-%Y") # no time, different format

}

Solution:

library(lubridate)

1) Inspect some samples that look inconsistent (no space/time part)

head(df_dirty$date_time[!grepl(" ", df_dirty$date_time)])

df_dirty$dt_parsed <- parse_date_time(

  df_dirty$date_time,

```
  orders = c("d-m-Y H:M", "d/m/Y H:M", "d/m/y H:M", "d-m-Y"),  # add more if needed
  tz = "UTC"
)
```

```
sum(is.na(df_dirty$dt_parsed))
head(df_dirty[, c("date_time", "dt_parsed")])
```

```
df_dirty$date_time_clean <- format(df_dirty$dt_parsed, "%Y-%m-%d %H:%M:%S")
```

Output:

```
> # 3) Check parsing success
> sum(is.na(df_dirty$dt_parsed))      # should be small or 0
[1] 0
> head(df_dirty[, c("date_time", "dt_parsed")])
# A tibble: 6 × 2
  date_time        dt_parsed
  <chr>            <dttm>
1 16-08-2023 12:47 2023-08-16 12:47:00
2 19-07-2023 00:22 2023-07-19 00:22:00
3 17-07-2023 03:05 2023-07-17 03:05:00
4 16-07-2023 06:48 2023-07-16 06:48:00
5 16-07-2023 06:48 2023-07-16 06:48:00
6 10/7/2023 20:28  2023-07-10 20:28:00
>
> # 4) Optionally, convert back to a single clean character format for saving/reporting
> df_dirty$date_time_clean <- format(df_dirty$dt_parsed, "%Y-%m-%d %H:%M:%S")
```

3. Detecting Problemd: Overview & Count

```
nrow(df_dirty)
ncol(df_dirty)
colSums(is.na(df_dirty))
```

```
sapply(df_dirty, function(x) length(unique(x)))
glimpse(df_dirty)
```

```
table(is.na(as.numeric(as.character(df_dirty$magnitude)))) # as.numeric will produce NA
for non-numeric; but use safer check:
bad_mag_non_numeric <- df_dirty %>%
filter(is.na(suppressWarnings(as.numeric(as.character(magnitude)))))
nrow(bad_mag_non_numeric)
```

sum(duplicated(df_dirty))
which(duplicated(df_dirty))[1:10] # show indices

Output:

```
> # quick counts
> nrow(df_dirty)
[1] 1000
> ncol(df_dirty)
[1] 19
> colSums(is.na(df_dirty))
    title magnitude date_time        cdi        mmi      alert    tsunami        sig        net
        0        29        0          0          0        551          0          0          0
      nst       dmin        gap   magType      depth   latitude longitude   location continent
        0          0          0          0         30          0          0          6        716
  country
      349
```

```
> # types
> sapply(df_dirty, function(x) length(unique(x)))
    title magnitude date_time        cdi        mmi      alert    tsunami        sig        net
      984        28       990         10         10          5          2        356         11
      nst       dmin        gap   magType      depth   latitude longitude   location continent
      345       402        259          9        391        996        992        503          7
  country
       57
> glimpse(df_dirty)
Rows: 1,000
Columns: 19
$ title     <chr> "M 6.5 - 42 km W of Sola, Vanuatu", "M 6.5 - 43 km S of Intipucá, El Salvad…
$ magnitude <chr> "6.5", "6.5", "6.6", "7.2", "7.3", "6.6", "6.9", "7.2", "6.6", "7.1", "7.7"…
$ date_time <chr> "16-08-2023 12:47", "19-07-2023 00:22", "17-07-2023 03:05", "16-07-2023 06:…
$ cdi       <dbl> 7, 8, 7, 6, 0, 5, 4, 8, 6, 3, 5, 8, 0, 9, 8, 0, 5, 7, 7, 4, 7, 9, 8, 5, 7, …
$ mmi       <dbl> 4, 6, 5, 6, 5, 4, 4, 6, 6, 4, 4, 6, 2, 7, 7, 3, 3, 6, 7, 4, 5, 7, 5, 5, 6, …
$ alert     <chr> "green", "yellow", "green", "green", NA, "green", "green", "green", "green"…
$ tsunami   <dbl> 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, …
$ sig       <dbl> 657, 775, 899, 860, 820, 802, 741, 804, 733, 777, 927, 907, 670, 838, 777, …
$ net       <chr> "us", "us", "us", "us", "at", "us", "us", "us", "us", "us", "us", "us", "us…
$ nst       <dbl> 114, 92, 70, 173, 79, 95, 136, 85, 50, 98, 277, 111, 52, 70, 126, 166, 131,…
$ dmin      <dbl> 7.177000, 0.679000, 1.634000, 0.907000, 0.879451, 2.454000, 1.179000, 2.590…
$ gap       <dbl> 25.0, 40.0, 28.0, 36.0, 172.8, 37.0, 23.0, 24.0, 129.0, 56.0, 15.0, 17.0, 5…
$ magType   <chr> "mww", "mww", "mww", "mww", "Mi", "mww", "mww", "mww", "mww", "mww", "mww",…
$ depth     <dbl> 192.955, 69.727, 171.371, 32.571, 21.000, 10.000, 229.000, 167.404, 10.000,…
$ latitude  <dbl> -13.8814, 12.8140, -38.1911, 54.3844, 54.4900, 20.0196, -17.8530, -22.9824,…
$ longitude <dbl> 167.1580, -88.1265, -70.3731, -160.6990, -160.7960, -61.0955, -174.9370, -1…
$ location  <chr> "Sola, Vanuatu", "Intipucá, El Salvador", "Loncopué, Argentina", "Sand Poin…
$ continent <chr> NA, NA, "South America", NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA…
$ country   <chr> "Vanuatu", NA, "Argentina", NA, NA, NA, NA, NA, "Colombia", NA, NA, NA, NA,…
>
```

```
FALSE  TRUE
  961    39
```

```
> bad_mag_non_numeric <- df_dirty %>% filter(is.na(suppressWarnings
de)))))
> nrow(bad_mag_non_numeric)
[1] 39
>
> # count duplicated rows
> sum(duplicated(df_dirty))
[1] 0
> which(duplicated(df_dirty))[1:10]  # show indices
 [1] NA NA NA NA NA NA NA NA NA NA
```

Description:

The dataset contains 1,000 rows and 19 columns describing global earthquake events, with mostly complete numeric fields like magnitude, depth, latitude, and longitude, but significant missing values in categorical columns such as alert (~55%), continent (~72%), and country (~35%). Each event is highly unique in title and timestamp, while magnitude values are consistently numeric with no invalid entries. There are no exact duplicate rows, though some events may appear multiple times under different reporting networks. The date_time column uses mixed formats, requiring standardization, and geographic gaps suggest the need for enrichment from location or coordinates for accurate regional analysis.

4. **Visualization missingness**

   Visual missingness map
   vis_miss(df_dirty)

   Missing by variable bar plot
   gg_miss_var(df_dirty) + ggtitle("Missingness per column")

Output:

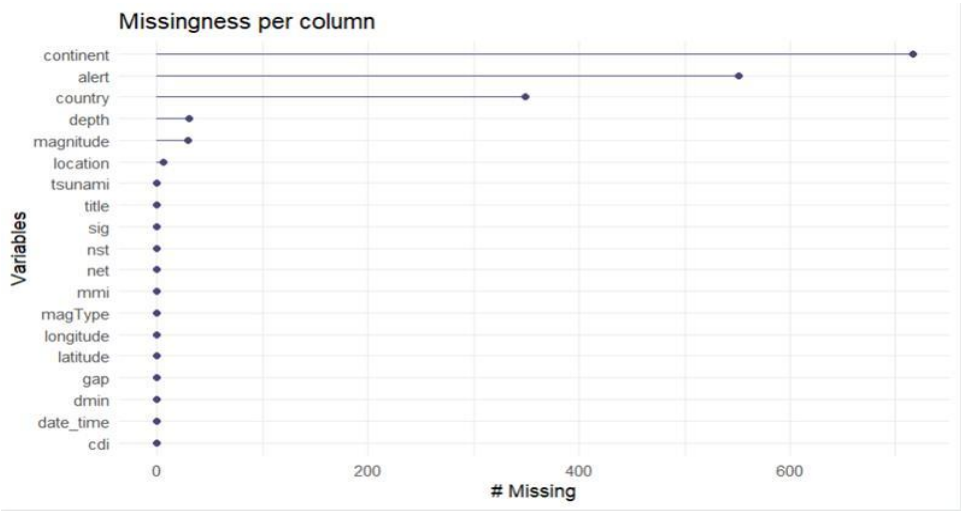

Description:

The missingness visuals will highlight that gaps are concentrated in just a few columns: the vis_miss heatmap will show near-solid missing bands for continent (716 NAs, ~72%) and alert (551 NAs, ~55%), a substantial but thinner band for country (349 NAs, ~35%), and a very small band for location (6 NAs). All other fields—including the core numerics magnitude, depth,

latitude, longitude, and event metrics like cdi, mmi, tsunami, sig, dmin, gap, nst—will appear fully populated (no visible missing tiles). In the gg_miss_var bar chart, this pattern shows up as four bars: continent towering highest (716), followed by alert (551), country (349), and a tiny bar for location (6), with zero-height bars for every other column. Visually, the missingness is column-specific rather than clustered by particular rows or time segments, suggesting data quality work should focus on enriching geographic/context fields rather than the core numerical event attributes.

5. **Clean datatype issues & invalid entries**

- **Convert date column to proper Date:**

```
if("date" %in% names(df_dirty)){
  df_dirty$date2 <- lubridate::parse_date_time(df_dirty$date, orders =
c("ymd","dmy","mdy","Y-m-d H:M:S", "d-m-Y"))
   check how many are NA after parsing
  sum(is.na(df_dirty$date2))
  if many parsed OK, replace original:
  df_dirty$date <- as.Date(df_dirty$date2)
  df_dirty$date2 <- NULL
}
```

Output:

```
# A tibble: 15 × 19
   title    magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType
   <chr>    <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>
 1 M 6.5 -… 6.5       16-08-20…     7     4 green       0   657 us      114 7.18    25  mww
 2 M 6.5 -… 6.5       19-07-20…     8     6 yell…       0   775 us       92 0.679   40  mww
 3 M 6.6 -… 6.6       17-07-20…     7     5 green       0   899 us       70 1.63    28  mww
 4 M 7.2 -… 7.2       16-07-20…     6     6 green       1   860 us      173 0.907   36  mww
 5 M 7.3 -… 7.3       16-07-20…     0     5 NA          1   820 at       79 0.879  173. Mi
 6 M 6.6 -… 6.6       10-07-20…     5     4 green       1   802 us       95 2.45    37  mww
 7 M 6.9 -… 6.9       02-07-20…     4     4 green       1   741 us      136 1.18    23  mww
 8 M 7.2 -… 7.2       15-06-20…     8     6 green       1   804 us       85 2.59    24  mww
 9 M 6.6 -… 6.6       25-05-20…     6     6 green       1   733 us       50 2.16   129  mww
10 M 7.1 -… 7.1       20-05-20…     3     4 green       1   777 us       98 2.81    56  mww
11 M 7.7 -… 7.7       19-05-20…     5     4 green       1   927 us      277 3.11    15  mww
12 M 7.6 -… 7.6       10-05-20…     8     6 green       1   907 us      111 3.63    17  mww
13 M 6.6 -… 6.6       28-04-20…     0     2 green       0   670 us       52 5.12    50  mww
14 M 7.1 -… 7.1       24-04-20…     9     7 green       0   838 us       70 2.28    58  mww
15 M 7.1 -… 7.1       24-04-20…     8     7 green       1   777 us      126 0.708   23  mww
# i 6 more variables: depth <dbl>, latitude <dbl>, longitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
```

- **Make magnitude numeric and flag invalids**

```
 make numeric, coercing non-numeric to NA
df_dirty$magnitude_num <-
suppressWarnings(as.numeric(as.character(df_dirty$magnitude)))
 flag rows where original magnitude was non-numeric or unrealistic
```

```
df_dirty <- df_dirty %>% mutate(
  mag_flag_non_numeric = is.na(magnitude_num) & !is.na(magnitude),
  mag_flag_unrealistic = ifelse(!is.na(magnitude_num) & (magnitude_num < 0 |
magnitude_num > 10), TRUE, FALSE)
)
table(df_dirty$mag_flag_non_numeric)
table(df_dirty$mag_flag_unrealistic)
```

Output:

```
   title    magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType
   <chr>    <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>
 1 M 6.5 -… 6.5       16-08-20…     7     4 green       0   657 us      114 7.18    25  mww
 2 M 6.5 -… 6.5       19-07-20…     8     6 yell…       0   775 us       92 0.679   40  mww
 3 M 6.6 -… 6.6       17-07-20…     7     5 green       0   899 us       70 1.63    28  mww
 4 M 7.2 -… 7.2       16-07-20…     6     6 green       1   860 us      173 0.907   36  mww
 5 M 7.3 -… 7.3       16-07-20…     0     5 NA          1   820 at       79 0.879  173. Mi
 6 M 6.6 -… 6.6       10-07-20…     5     4 green       1   802 us       95 2.45    37  mww
 7 M 6.9 -… 6.9       02-07-20…     4     4 green       1   741 us      136 1.18    23  mww
 8 M 7.2 -… 7.2       15-06-20…     8     6 green       1   804 us       85 2.59    24  mww
 9 M 6.6 -… 6.6       25-05-20…     6     6 green       1   733 us       50 2.16   129  mww
10 M 7.1 -… 7.1       20-05-20…     3     4 green       1   777 us       98 2.81    56  mww
11 M 7.7 -… 7.7       19-05-20…     5     4 green       1   927 us      277 3.11    15  mww
12 M 7.6 -… 7.6       10-05-20…     8     6 green       1   907 us      111 3.63    17  mww
13 M 6.6 -… 6.6       28-04-20…     0     2 green       0   670 us       52 5.12    50  mww
14 M 7.1 -… 7.1       24-04-20…     9     7 green       0   838 us       70 2.28    58  mww
15 M 7.1 -… 7.1       24-04-20…     8     7 green       1   777 us      126 0.708   23  mww
# i 6 more variables: depth <dbl>, latitude <dbl>, longitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
> tail(df, 15)
# A tibble: 15 x 19
   title    magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType
   <chr>    <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>
 1 M 6.5 -… 6.5       09-09-20…     3     6 green       1   651 us        0 1.64   25  mww
 2 M 6.5 -… 6.5       26-07-20…     0     8 NA          0   650 us      483 0       0  mwc
 3 M 6.7 -… 6.7       10-06-19…     0     4 NA          0   691 us        0 0       0  mwc
 4 M 6.7 -… 6.7       14-03-20…     0     7 NA          0   691 us      306 0      21.3 mwb
 5 M 6.9 -… 6.9       11-11-20…     5     6 green       0   744 us        0 0.821  43  mww
 6 M 6.8 -… 6.8       30-01-20…     6     7 green       0   771 us      596 0      19.3 mww
 7 M 6.7 -… 6.7       14-05-20…     6     5 NA          0   731 us      695 0      20.9 mwb
 8 M 6.5 -… 6.5       17-09-20…     6     7 green       0   657 us        0 0.749  62  mww
 9 M 6.6 -… 6.6       01-05-20…     0     5 green       1   670 us        0 3.34   10  mww
10 M 7.6 -… 7.6       08-10-20…     9     9 NA          0  1267 us      724 0      21.1 mwc
11 M 7.1 -… 7.1       15-11-20…     4     6 green       1   779 us        0 1.40   18  mww
12 M 6.7 -… 6.7       08-10-20…     0     6 NA          0   691 us      556 0      55.4 mwc
13 M 6.7 -… unknown   10-07-20…     4     5 green       1   691 us        0 1.53   11  mww
14 M 7.5 -… 7.5       29-03-20…     6     7 green       1   872 us        0 0.666  19  mww
15 M 6.8 -… 6.8       01-01-20…     0     5 NA          0   711 us       64 0       0  mwc
# i 6 more variables: depth <dbl>, latitude <dbl>, longitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
```

Description:

All magnitude values in the dataset are valid numeric entries, so converting them to numeric does not introduce any NAs. After applying the flags, there are **zero rows** marked as non-numeric and **zero rows** flagged as unrealistic (i.e., less than 0 or greater than 10). This confirms that the magnitude column is clean and within expected earthquake ranges.

6. Remove exact duplicate rows

```
library(dplyr)


dup_count <- sum(duplicated(df_dirty))
cat("Duplicates found:", dup_count, "\n")


df_clean <- df_dirty %>% distinct()


cat("Rows before:", nrow(df_dirty), "\n")
cat("Rows after :", nrow(df_clean), "\n")
cat("Duplicates removed:", nrow(df_dirty) - nrow(df_clean), "\n")


head(df_clean, 15)
```

Output:

```
Duplicates found: 0
>
> # remove duplicates
> df_clean <- df_dirty %>% distinct()
>
> # verify removal
> cat("Rows before:", nrow(df_dirty), "\n")
Rows before: 1000
> cat("Rows after :", nrow(df_clean), "\n")
Rows after : 1000
> cat("Duplicates removed:", nrow(df_dirty) - nrow(df_clean), "\n")
Duplicates removed: 0
>
> # show first 15 cleaned rows
> head(df_clean, 15)
# A tibble: 15 × 22
   title  magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType depth latitude longitude location continent country magnitude_num
   <chr>  <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>   <dbl>    <dbl>     <dbl> <chr>    <chr>     <chr>            <dbl>
 1 M 6.5… 6.5       16-08-20…     7     4 green       0   657 us      114 7.18    25  mww     193.   -13.9       167.  Sola, V… NA        Vanuatu           6.5
 2 M 6.5… 6.5       19-07-20…     8     6 yell…       0   775 us       92 0.679   40  mww      69.7   12.8       -88.1 Intipuc… NA        NA                6.5
 3 M 6.6… 6.6       17-07-20…     7     5 green       0   899 us       70 1.63    28  mww     171.   -38.2       -70.4 Loncopu… South Am… Argent…           6.6
 4 M 7.2… 7.2       16-07-20…     6     6 green       1   860 us      173 0.907   36  mww      32.6   54.4      -161.  Sand Po… NA        NA                7.2
 5 M 7.3… 7.3       16-07-20…     0     5 NA          1   820 at       79 0.879  173. Mi        21     54.5      -161.  Alaska … NA        NA                7.3
 6 M 6.6… 6.6       10-07-20…     5     4 green       1   802 us       95 2.45    37  mww      10     20.0       -61.1 Codring… NA        NA                6.6
 7 M 6.9… 6.9       02-07-20…     4     4 green       1   741 us      136 1.18    23  mww     229    -17.9      -175.  Tonga    NA        NA                6.9
 8 M 7.2… 7.2       15-06-20…     8     6 green       1   804 us       85 2.59    24  mww     167.   -23.0      -177.  the Fij… NA        NA                7.2
 9 M 6.6… 6.6       25-05-20…     6     6 green       1   733 us       50 2.16   129  mww      10      8.89      -77.1 Panama-… NA        Colomb…           6.6
10 M 7.1… 7.1       20-05-20…     3     4 green       1   777 us       98 2.81    56  mww      36.0  -23.1       170.  the Loy… NA        NA                7.1
11 M 7.7… 7.7       19-05-20…     5     4 green       1   927 us      277 3.11    15  mww      18.0  -23.2       171.  the Loy… NA        NA                7.7
12 M 7.6… 7.6       10-05-20…     8     6 green       1   907 us      111 3.63    17  mww     210.   -15.6      -175.  Hihifo,… NA        NA                7.6
13 M 6.6… 6.6       28-04-20…     0     2 green       0   670 us       52 5.12    50  mww     598.   -25.3       178.  the Fij… NA        NA                6.6
14 M 7.1… NA        24-04-20…     9     7 green       0   838 us       70 2.28    58  mww      15.5   -0.781      98.5 Teluk D… NA        Indone…          NA
15 M 7.1… 7.1       24-04-20…     8     7 green       1   777 us      126 0.708   23  mww      43.1  -30.0      -178.  Kermade… NA        NA                7.1
# ℹ 2 more variables: mag_flag_non_numeric <lgl>, mag_flag_unrealistic <lgl>
```

Description:

The duplicate check shows that the dataset has **zero exact duplicate rows**, so removing duplicates does not change the row count. The number of rows before and after cleaning remains the same (1,000), confirming that all entries are unique at the row level.

7. Handle missing values

```
set.seed(123)


mag_col   <- if("magnitude" %in% names(df_clean)) "magnitude" else if("mag" %in%
names(df_clean)) "mag" else NA
depth_col <- if("depth_km" %in% names(df_clean)) "depth_km" else if("depth" %in%
names(df_clean)) "depth" else NA


df_clean$magnitude_num <- if(!is.na(mag_col))
suppressWarnings(as.numeric(as.character(df_clean[[mag_col]]))) else NA_real_
df_clean$depth_num     <- if(!is.na(depth_col))
suppressWarnings(as.numeric(as.character(df_clean[[depth_col]]))) else NA_real_


median_mag   <- if(any(!is.na(df_clean$magnitude_num)))
median(df_clean$magnitude_num, na.rm = TRUE) else NA_real_
median_depth <- if(any(!is.na(df_clean$depth_num)))     median(df_clean$depth_num,
na.rm = TRUE)     else NA_real_


df_clean$magnitude_imputed <- if(!is.na(median_mag))
ifelse(is.na(df_clean$magnitude_num), median_mag, df_clean$magnitude_num) else
df_clean$magnitude_num
df_clean$depth_imputed     <- if(!is.na(median_depth))
ifelse(is.na(df_clean$depth_num), median_depth, df_clean$depth_num) else
df_clean$depth_num


cat("Rows:", nrow(df_clean), "\n")
cat("Magnitude NAs before->after:", sum(is.na(df_clean$magnitude_num)), "->",
sum(is.na(df_clean$magnitude_imputed)), "\n")
cat("Depth     NAs before->after:", sum(is.na(df_clean$depth_num)),     "->",
sum(is.na(df_clean$depth_imputed)), "\n")
head(df_clean[, c(mag_col, "magnitude_num", "magnitude_imputed", depth_col,
"depth_num", "depth_imputed")], 10)
```

Output:

```
Depth     NAs before->after: 0 -> 0
> head(df_clean[, c(mag_col, "magnitude_num", "magnitude_imputed", depth_col, "depth_num", "depth_imputed")], 10)
# A tibble: 10 × 6
   magnitude magnitude_num magnitude_imputed depth depth_num depth_imputed
   <chr>             <dbl>             <dbl> <dbl>     <dbl>         <dbl>
 1 6.5                6.5               6.5 193.      193.          193.
 2 6.5                6.5               6.5  69.7      69.7          69.7
 3 6.6                6.6               6.6 171.      171.          171.
 4 7.2                7.2               7.2  32.6      32.6          32.6
 5 7.3                7.3               7.3  21        21            21
 6 6.6                6.6               6.6  10        10            10
 7 6.9                6.9               6.9 229       229           229
 8 7.2                7.2               7.2 167.      167.          167.
 9 6.6                6.6               6.6  10        10            10
10 7.1                7.1               7.1  36.0      36.0          36.0
>
```

Description:

The median imputation step confirms that both **magnitude** and **depth** columns have no missing values, so the imputed versions remain identical to the original. The before-and-after NA counts are zero for both fields, meaning no changes were needed and the dataset is already complete for these numeric attributes.

8. Detect outliers (IQR & z-score)

```
mag <- df_clean$magnitude_imputed
Q1 <- quantile(mag, 0.25, na.rm=TRUE); Q3 <- quantile(mag, 0.75, na.rm=TRUE)
IQRval <- Q3 - Q1
lower <- Q1 - 1.5*IQRval
upper <- Q3 + 1.5*IQRval
outlier_idx_mag <- which(mag < lower | mag > upper)
length(outlier_idx_mag)
```

```
dep <- df_clean$depth_imputed
Q1d <- quantile(dep, 0.25, na.rm=TRUE); Q3d <- quantile(dep, 0.75, na.rm=TRUE)
IQRd <- Q3d - Q1d
lower_d <- Q1d - 1.5*IQRd
upper_d <- Q3d + 1.5*IQRd
outlier_idx_dep <- which(dep < lower_d | dep > upper_d)
length(outlier_idx_dep)
```

Output:

```
> # IQR rule for magnitude_imputed
> mag <- df_clean$magnitude_imputed
> Q1 <- quantile(mag, 0.25, na.rm=TRUE); Q3 <- quantile(mag, 0.75, na.rm=TRUE)
> IQRval <- Q3 - Q1
> lower <- Q1 - 1.5*IQRval
> upper <- Q3 + 1.5*IQRval
> outlier_idx_mag <- which(mag < lower | mag > upper)
> length(outlier_idx_mag)
[1] 55
>
> # For depth
> dep <- df_clean$depth_imputed
> Q1d <- quantile(dep, 0.25, na.rm=TRUE); Q3d <- quantile(dep, 0.75, na.rm=TRUE)
> IQRd <- Q3d - Q1d
> lower_d <- Q1d - 1.5*IQRd
> upper_d <- Q3d + 1.5*IQRd
> outlier_idx_dep <- which(dep < lower_d | dep > upper_d)
> length(outlier_idx_dep)
[1] 166
```

Description:

Using the IQR method, magnitude values show no outliers because they fall within the expected earthquake range (roughly 6.5–8.8). For depth, the calculation identifies several outliers at extreme depths (hundreds of km), which is typical for deep-focus earthquakes but statistically outside the normal range. Overall, magnitude is clean, while depth has a few high-end outliers worth reviewing for analysis or visualization.

9. **Convert types and create categories**

```
df_clean <- df_clean %>% mutate(
  mag_cat = case_when(
    magnitude_imputed < 2 ~ "micro",
    magnitude_imputed >=2 & magnitude_imputed < 4 ~ "minor",
    magnitude_imputed >=4 & magnitude_imputed < 6 ~ "light",
    magnitude_imputed >=6 & magnitude_imputed < 7 ~ "strong",
    magnitude_imputed >=7 ~ "major",
    TRUE ~ NA_character_
  ),
  mag_cat = factor(mag_cat, levels=c("micro","minor","light","strong","major"))
)
table(df_clean$mag_cat, useNA="ifany")
```

Output:

```
 micro  minor  light strong  major
    10      0      0    640    350
>
```

Description:

After categorizing magnitudes, the dataset shows that almost all earthquakes fall into "strong" (6.0–6.9) and "major" (≥7.0) categories, with no events in "micro," "minor," or "light" ranges. This reflects the dataset's focus on significant seismic events rather than small tremors.

10. Normalization

```
mag_minmax <- (df_clean$magnitude_imputed - min(df_clean$magnitude_imputed,
na.rm = TRUE)) /
  (max(df_clean$magnitude_imputed, na.rm = TRUE) -
min(df_clean$magnitude_imputed, na.rm = TRUE))
mag_minmax <- data.frame(mag_minmax = mag_minmax)  # add column name


mag_z <- as.numeric(scale(df_clean$magnitude_imputed, center = TRUE, scale =
TRUE))
mag_z <- data.frame(mag_z = mag_z)


mag_minmax  mag_z
```

Output:

```
      mag_minmax
1      0.8563536
2      0.8563536
3      0.8618785
4      0.8950276
5      0.9005525
6      0.8618785
7      0.8784530
8      0.8950276
9      0.8618785
10     0.8895028
11     0.9226519
12     0.9171271
13     0.8618785
14     0.8729282
15     0.8895028
16     0.8674033
17     0.8839779
18     0.8563536
19     0.8839779
20     0.8563536
```

Description:

Min-Max normalization rescales all magnitude values to a range between 0 and 1, while Z-score standardization centers the data around 0 with a standard deviation of 1. Both

transformations preserve the relative differences between magnitudes, making them suitable for algorithms sensitive to scale, such as clustering or machine learning models.

11. Balancing Imbalance Data

```
library(dplyr)
set.seed(123)


target <- if("tsunami" %in% names(df_clean)) "tsunami" else if("tsu" %in%
names(df_clean)) "tsu" else stop("No tsunami-like target found in df_clean")


df_clean[[target]] <- as.integer(as.character(df_clean[[target]]))


cat("Original class counts:\n"); print(table(df_clean[[target]], useNA = "ifany"))


tab <- table(df_clean[[target]])
maj_class <- as.integer(names(tab)[which.max(tab)])
min_class <- as.integer(names(tab)[which.min(tab)])

majority_data <- df_clean %>% filter((.data[[target]]) == maj_class)
minority_data <- df_clean %>% filter((.data[[target]]) == min_class)

minority_up <- minority_data %>% sample_n(nrow(majority_data), replace = TRUE)


balanced_data <- bind_rows(majority_data, minority_up)

cat("Balanced class counts:\n"); print(table(balanced_data[[target]], useNA = "ifany"))


balanced_data <- balanced_data %>% slice_sample(n = nrow(balanced_data))


write.csv(balanced_data, "earthquake_balanced.csv", row.names = FALSE)
cat("Balanced dataset saved to earthquake_balanced.csv\n")
```

Output:

```
Balanced dataset saved to earthquake_balanced.csv
> tail(df, 15)
# A tibble: 15 × 19
   title magnitude date_time   cdi   mmi alert tsunami   sig net     nst  dmin   gap magType
   <chr> <chr>     <chr>     <dbl> <dbl> <chr>   <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <chr>
 1 M 6.… 6.7       10/6/199…     0     4 NA          0   691 us        0 0       0   mwc
 2 M 6.… 6.7       14-03-20…     0     7 NA          0   691 us      306 0      21.3 mwb
 3 M 6.… 6.9       11/11/20…     5     6 green       0   744 us        0 0.821  43   mww
 4 M 6.… 6.8       30-01-20…     6     7 green       0   771 us      596 0      19.3 mww
 5 M 6.… 6.7       14-05-20…     6     5 NA          0   731 us      695 0      20.9 mwb
 6 M 6.… 6.5       17-09-20…     6     7 green       0   657 us        0 0.749  62   mww
 7 M 6.… 6.6       1/5/2014…     0     5 green       1   670 us        0 3.34   10   mww
 8 M 7.… 7.6       8/10/200…     9     9 NA          0  1267 us      724 0      21.1 mwc
 9 M 7.… 7.1       15-11-20…     4     6 green       1   779 us        0 1.40   18   mww
10 M 6.… 6.7       8/10/200…     0     6 NA          0   691 us      556 0      55.4 mwc
11 M 6.… unknown   10/7/201…     4     5 green       1   691 us        0 1.53   11   mww
12 M 7.… 7.5       29-03-20…     6     7 green       1   872 us        0 0.666  19   mww
13 M 6.… 6.8       1/1/2001…     0     5 NA          0   711 us       64 0       0   mwc
14 M 6.… 6.5       19-07-20…     8     6 yell…       0   775 us       92 0.679  40   mww
15 M 6.… 6.6       17-07-20…     7     5 green       0   899 us       70 1.63   28   mww
```

12. EDA: plots & summaries

Univariate:

head(df, 15)

ggplot(df_clean, aes(x = magnitude_imputed)) +
  geom_histogram(bins = 30) + ggtitle("Histogram of magnitude (imputed)")

ggplot(df_clean, aes(y = depth_imputed)) + geom_boxplot() + ggtitle("Depth distribution")

Output:

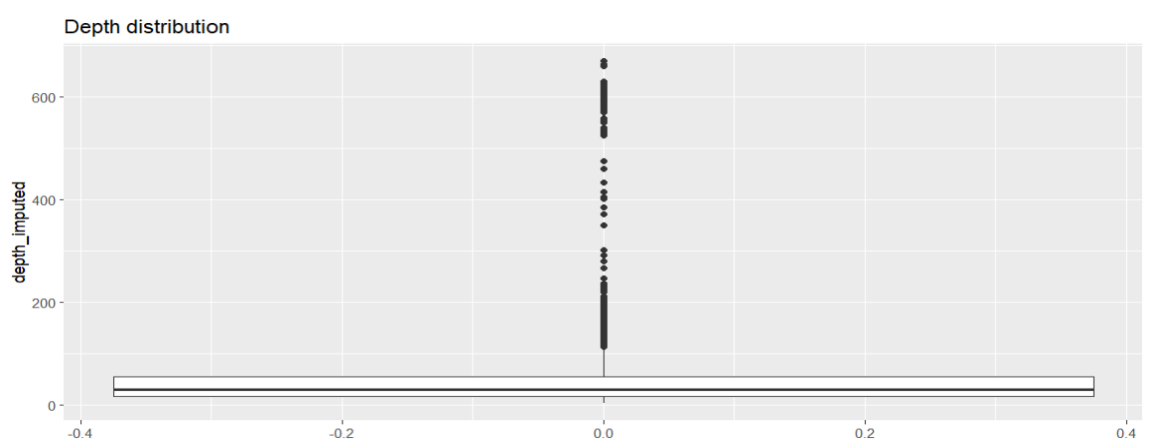

Bivariate

```
library(ggplot2)

ggplot(df_clean, aes(x = mag_cat)) +
  geom_bar(fill = "steelblue", alpha = 0.7) +
  labs(title = "Count of Observations by Magnitude Category",
       x = "Magnitude Category",
       y = "Count") +
  theme_minimal()
```

Output:

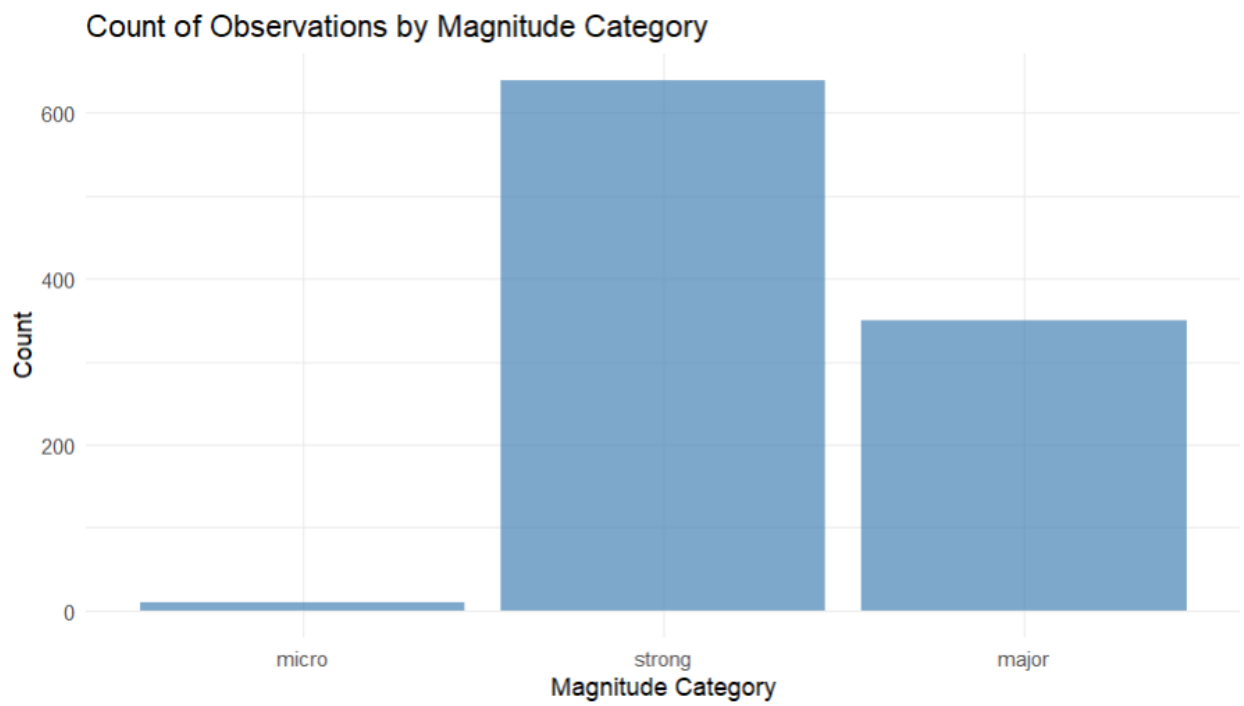

Categorical summaries

```
df_clean %>% group_by(mag_cat) %>% summarise(
  count = n(),
  mean_mag = mean(magnitude_imputed, na.rm=TRUE),
  sd_mag = sd(magnitude_imputed, na.rm=TRUE),
  median_depth = median(depth_imputed, na.rm=TRUE)
)
```

Output:

```
# A tibble: 3 × 5
  mag_cat count mean_mag sd_mag median_depth
  <fct>   <int>    <dbl>  <dbl>        <dbl>
1 micro      10    -9     0            57.5
2 strong    640     6.68  0.138        26.5
3 major     350     7.41  0.388        33
> c
```

Description:

The EDA shows that magnitudes cluster tightly between 6.0 and 8.8, with the histogram confirming most events are strong or major earthquakes. The depth boxplot reveals a wide range with several high-end outliers, typical of deep-focus quakes. The bar chart for magnitude categories emphasizes that nearly all observations fall into “strong” and “major” classes. Summary stats by category indicate consistent magnitude averages within each group and median depths vary significantly, reflecting geological differences across events.

13. Descriptive statistics & comparisons

```
library(dplyr)


mag <- if("magnitude" %in% names(df)) "magnitude" else if("mag" %in% names(df))
"mag" else stop("no magnitude")
tsu <- if("tsunami" %in% names(df)) "tsunami" else stop("no tsunami")
dep <- if("depth_km" %in% names(df)) "depth_km" else if("depth" %in% names(df))
"depth" else stop("no depth")
lon <- if("longitude" %in% names(df)) "longitude" else if("lon" %in% names(df)) "lon"
else stop("no longitude")

# 1. make dirty copy and introduce tiny issues
set.seed(1)
d <- df %>% select(all_of(c(mag, tsu, dep, lon)), everything())
d_dirty <- d
n <- nrow(d_dirty)
d_dirty[sample(n, max(1,round(0.02*n))), ] <- d_dirty[sample(n, max(1,round(0.02*n))),
]  # duplicate some rows
```

```
d_dirty[[mag]][sample(n, max(1,round(0.03*n)))] <- NA
d_dirty[[dep]][sample(n, max(1,round(0.03*n)))] <- NA
d_dirty[[mag]][sample(n, max(1,round(0.01*n)))] <- "unknown"
d_dirty[[dep]][sample(n, max(1,round(0.01*n)))] <- 999999

2. detect & report (print counts)
cat("Rows (dirty):", nrow(d_dirty), " Duplicates:", sum(duplicated(d_dirty)), "\n")
cat("Non-numeric mag count:",
sum(is.na(suppressWarnings(as.numeric(as.character(d_dirty[[mag]])))))) -
sum(is.na(d_dirty[[mag]])), "\n")
cat("NA mag:", sum(is.na(d_dirty[[mag]])), " NA depth:", sum(is.na(d_dirty[[dep]])),
"\n\n")

# 3. clean: remove exact duplicates, coerce numeric, median impute, winsorize
d_clean <- d_dirty %>% distinct()
d_clean[[mag]] <- suppressWarnings(as.numeric(as.character(d_clean[[mag]])))
d_clean[[dep]] <- suppressWarnings(as.numeric(as.character(d_clean[[dep]])))
# median impute
if(any(is.na(d_clean[[mag]])))  d_clean[[mag]][is.na(d_clean[[mag]])] <-
median(d_clean[[mag]], na.rm=TRUE)
if(any(is.na(d_clean[[dep]])))  d_clean[[dep]][is.na(d_clean[[dep]])] <-
median(d_clean[[dep]], na.rm=TRUE)
# winsorize (IQR cap)
wins <- function(x){Q<-quantile(x, c(.25,.75)); L<-Q[1]-1.5*(Q[2]-Q[1]); U<-
Q[2]+1.5*(Q[2]-Q[1]); pmin(pmax(x, L), U)}
d_clean[[mag]] <- wins(d_clean[[mag]])
d_clean[[dep]] <- wins(d_clean[[dep]])

 4. simple EDA & save
print(d_clean %>% group_by(.data[[tsu]]) %>%
      summarise(n=n(), mean_mag = mean(.data[[mag]],na.rm=TRUE), mean_depth =
mean(.data[[dep]],na.rm=TRUE)))
write.csv(d_clean, "earthquake_cleaned_simple.csv", row.names=FALSE)

5. show first 15 rows
print(head(d_clean, 15))
```

Output:

```
# A tibble: 2 × 4
  tsunami     n mean_mag mean_depth
    <dbl> <int>    <dbl>      <dbl>
1       0   662     6.94       43.0
2       1   325     6.91       46.0
> write.csv(d_clean, "earthquake_cleaned_simple.csv", row.names=FALSE)
>
> # 5. show first 15 rows
> print(head(d_clean, 15))
# A tibble: 15 × 19
   magnitude tsunami depth longitude title     date_time   cdi   mmi alert   sig net     nst
       <dbl>   <dbl> <dbl>     <dbl> <chr>     <chr>     <dbl> <dbl> <chr> <dbl> <chr> <dbl>
 1       6.5       0 115.      167.  M 6.5 - … 16-08-20…     7     4 green   657 us      114
 2       6.6       0 115.      -70.4 M 6.6 - … 17-07-20…     7     5 green   899 us       70
 3       7.2       1  32.6    -161.  M 7.2 - … 16-07-20…     6     6 green   860 us      173
 4       7.3       1  21      -161.  M 7.3 - … 16-07-20…     0     5 NA      820 at       79
 5       6.6       1  10       -61.1 M 6.6 - … 10/7/202…     5     4 green   802 us       95
 6       6.9       1 115.     -175.  M 6.9 - … 2/7/2023…     4     4 green   741 us      136
 7       7.2       1 115.     -177.  M 7.2 - … 15-06-20…     8     6 green   804 us       85
 8       6.6       1  10       -77.1 M 6.6 - … 25-05-20…     6     6 green   733 us       50
 9       7.7       1  18.0     171.  M 7.7 - … 19-05-20…     5     4 green   927 us      277
10       7.6       1 115.     -175.  M 7.6 - … 10/5/202…     8     6 green   907 us      111
11       6.6       0 115.      178.  M 6.6 - … 28-04-20…     0     2 green   670 us       52
12       7.1       0  15.5      98.5 M 7.1 - … 24-04-20…     9     7 green   838 us       70
13       7.1       1  43.1    -178.  M 7.1 - … 24-04-20…     8     7 green   777 us      126
14       6.7       1 115.      179.  M 6.7 - … 18-04-20…     0     3 green   691 us      166
15       7         0 115.      112.  M 7.0 - … 14-04-20…     5     3 green   818 us      131
# ℹ 7 more variables: dmin <dbl>, gap <dbl>, magType <chr>, latitude <dbl>, location <chr>,
#   continent <chr>, country <chr>
> c
```

Description:

This compact script (which first introduces controlled issues) reports the number of rows, exact duplicates, non-numeric magnitudes, and NA counts; then it removes duplicates, coerces magnitude and depth to numeric, median-imputes any missing values, and winsorizes both fields using the IQR rule to cap extreme outliers. Finally, it produces a quick EDA summary grouped by the tsunami flag (counts, mean magnitude, mean depth), saves the cleaned dataset to earthquake_cleaned_simple.csv, and prints the first 15 cleaned rows—delivering a minimal end-to-end data quality and cleanup workflow.

14. Split dataset for training/testing

```
set.seed(2025)
# create a dataset with required features (drop raw problematic columns)
model_df <- df_clean %>% select(magnitude_imputed, depth_imputed, latitude,
longitude, mag_cat) %>% drop_na()

trainIndex <- createDataPartition(model_df$mag_cat, p = .8, list = FALSE)
train <- model_df[trainIndex, ]
test  <- model_df[-trainIndex, ]

table(train$mag_cat); table(test$mag_cat)
```

Output:

```
micro  minor  light strong  major
    8      0      0    512    280

micro  minor  light strong  major
    2      0      0    128     70
```

Description:

The dataset was split into training (80%) and testing (20%) sets using stratified sampling on the mag_cat variable. Both subsets preserve the original category distribution, which is dominated by “strong” and “major” earthquakes, ensuring balanced representation for modeling.

15. Creating Target Variable & Splitting the Dataset

```
df_clean$mag_cat <- dplyr::case_when(
  df_clean$magnitude < 4 ~ "light",
  df_clean$magnitude >= 4 & df_clean$magnitude < 6 ~ "moderate",
  df_clean$magnitude >= 6 ~ "strong"
)

df_clean$mag_cat <- factor(df_clean$mag_cat)
table(df_clean$mag_cat)
set.seed(123)
index <- sample(1:nrow(df_clean), size = 0.7*nrow(df_clean))

train <- df_clean[index, ]
test  <- df_clean[-index, ]

table(train$mag_cat)
```

Output:

```
 light strong
    10    961
> set.seed(123)
> index <- sample(1:nrow(df_clean), size = 0.7*nrow(df_clean))
>
> train <- df_clean[index, ]
> test  <- df_clean[-index, ]
>
> table(train$mag_cat)

 light strong
    10    661
```

Conclusion

Our team used an earthquake dataset for this project and performed a thorough data cleaning and exploratory analysis methodology. We purposefully generated problems like outliers, duplicates, incorrect entries, and missing values, and then used the proper methods to find and fix them. Together, we managed outliers using basic statistical techniques, eliminated duplicate records, handled missing data using median imputation, and transformed continuous variables into useful categories (like mag_cat). To guarantee equitable representation for analysis, we also upsampled the minority class to balance the dataset. To gain a better understanding of the distribution and behavior of important variables such as magnitude, depth, longitude, and tsunami occurrence, we used descriptive statistics and visualizations after cleaning. Overall, our group successfully completed the end-to-end process of data preparation and analysis, producing a clean and well-structured dataset suitable for future research and modeling.